

Data Analysis with Python

Notes of IBM Data Science Professional Certificate Courses on Coursera

[View on GitHub](#)

Data Analysis with Python

Data Analysis with Python	1
Data Analysis with Python	1
Datasets	2
Exporting to different formats in Python.....	2
Preprocessing Data in Python.....	3
How to deal with missing data	3
Data Formatting in Python	4
Data Normalization in Python	4
Binning	4
Turning categorical variables into quantitative variables in Python.....	5
Exploratory Data Analysis (EDA).....	6
Descriptive Statistics - Describe()	6
Grouping data	6
Correlation.....	8
Correlation - Statistics	9
Association between two categorical variables: Chi-Square	10
Model Development	12
Linear Regression and Multiple Linear Regression	12
Model Evaluation using Visualization	13
Polynomial Regression and Pipelines.....	15
Measures for In-Sample Evaluation	18
Model Evaluation and Refinement	19
Function <code>cross_val_score()</code>	20
Function <code>cross_val_predict()</code>	20
Overfitting, Underfitting and Model Selection.....	20
Ridge Regression.....	21
Grid Search.....	22

Datasets

Understanding Datasets

Each of the attributes in the dataset

No.	Attribute name	attribute range	No.	Attribute name	attribute range
1	symboling	-3, -2, -1, 0, 1, 2, 3.	14	curb-weight	continuous from 1488 to 4066.
2	normalized-losses	continuous from 65 to 256.	15	engine-type	dohc, dohc, l, ohc, ohcf, ohcv, rotor.
3	make	audi, bmw, etc.	16	num-of-cylinders	eight, five, four, six, three, twelve, two.
4	fuel-type	diesel, gas.	17	engine-size	continuous from 61 to 326.
5	aspiration	std, turbo.	18	fuel-system	1bbl, 2bbl, 4bbl, idi, mfi, mpfi, spdi, spfi.
6	num-of-doors	four, two.	19	bore	continuous from 2.54 to 3.94.
7	body-style	hardtop, wagon, etc.	20	stroke	continuous from 2.07 to 4.17.
8	drive-wheels	4wd, fwd, rwd.	21	compression-ratio	continuous from 7 to 23.
9	engine-location	front, rear.	22	horsepower	continuous from 48 to 288.
10	wheel-base	continuous from 86.6 to 120.9.	23	peak-rpm	continuous from 4150 to 6600.
11	length	continuous from 141.1 to 208.1.	24	city-mpg	continuous from 13 to 49.
12	width	continuous from 60.3 to 72.3.	25	highway-mpg	continuous from 16 to 54.
13	height	continuous from 47.8 to 59.8.	26	price	continuous from 5118 to 45400.

Data source: <https://archive.ics.uci.edu/ml/machine-learning-databases/autos/>

Importing Data

- Process of loading and reading data into Python from various resources.
- **Two important properties:**
 - **Format**
 - various formats: .csv, .json, .xlsx, .hdf
 - **File Path of dataset**
 - Computer: /Desktop/mydata.csv
 - Internet: <https://archive.ics.uci.edu/autos/imports-85.data>

Exporting to different formats in Python

Data Format	Read	Save
csv	<code>pd.read_csv()</code>	<code>df.to_csv()</code>
json	<code>pd.read_json()</code>	<code>df.to_json()</code>
Excel	<code>pd.read_excel()</code>	<code>df.to_excel()</code>
sql	<code>pd.read_sql()</code>	<code>df.to_sql()</code>

Basic insights from the data

- Understand your data before you begin any analysis
- Should check:
 - data types
 - `df.dtypes`
 - data distribution
 - `df.describe()`
 - `df.describe(include="all")`, provides full summary statistics
 - `unique`
 - `top`
 - `freq`
- Locate potential issues with the data
 - potential info and type mismatch
 - compatibility with python methods

Writing code using DB-API

```
from dbmodule import connect

#Create connection object
connection = connect('databasename', 'username', 'pswd')

#Create a cursor object
cursor = connection.cursor()

#Run queries
cursor.execute('select * from mytable')
results = cursor.fetchall()

#Free resources
Cursor.close()
connection.close()
```

Preprocessing Data in Python

- Identify and handle missing values
- Data formatting
- Data normalization (centering / scaling)
- Data binning
- Turning categorical values to numeric variables

How to deal with missing data

- Check with the data collection source
- Drop the missing values
 - drop the variable
 - drop the data entry
- Replace the missing values

- replace it with an average (of similar datapoints)
 - replace it by frequency
 - replace it based on other functions
- Leave it as missing data

```
df.dropna(subset=["price"], axis=0, inplace=True)
```

is equivalent to

```
df = df.dropna(subset=["price"], axis=0)
```

Data Formatting in Python

Non-formatted:

- confusing
- hard to aggregate
- hard to compare

Formatted:

- more clear
- easy to aggregate
- easy to compare

Correcting data types

- use `df.dtypes()` to identify data type
- use `df.astype()` to convert data type
 - e.g. `df["price"] = df["price"].astype("int")`

Data Normalization in Python

Approaches for normalization:

- Simple feature scaling: $x_{\text{new}} = x_{\text{old}}/x_{\text{max}}$
 - `df["length"] = df["length"] / df["length"].max()`
- Min-Max: $x_{\text{new}} = (x_{\text{old}} - x_{\text{min}})/(x_{\text{max}} - x_{\text{min}})$
 - `df["length"] = (df["length"] - df["length"].min()) / (df["length"].max() - df["length"].min())`
- Z-score: $x_{\text{new}} = (x_{\text{old}} - \mu)/\sigma$
 - `df["length"] = (df["length"] - df["length"].mean()) / df["length"].std()`

Binning

Binning

- Binning: Grouping of values into "bins"
- Converts numeric into categorical variables
- Group a set of numerical values into a set of "bins"
- "price" is a feature range from 5,000 to 45,500 (in order to have a **better representation** of price)

price: 5000, 10000, 12000, 12000, 30000, 31000, 39000, 44000, 44500

bins: 

```
bins = np.linspace(min(df["price"]), max(df["price"]), 4)
group_names = ["Low", "Medium", "High"]
df["price-binned"] = pd.cut(df["price"], bins, labels=group_names, include_lowest=True)
```

Turning categorical variables into quantitative variables in Python

Categorical → Numeric

Solution:

- Add dummy variables for each unique category
- Assign 0 or 1 in each category

Car	Fuel	...	gas	diesel
A	gas	...	1	0
B	diesel	...	0	1
C	gas	...	1	0
D	gas	...	1	0

"One-hot encoding"

Dummy variables in Python pandas

- Use pandas.get_dummies() method.
- Convert categorical variables to dummy variables (0 or 1)

fuel		gas	diesel
gas	→	1	0
diesel		0	1
gas		1	0
gas		1	0

```
pd.get_dummies(df['fuel'])
```

Exploratory Data Analysis (EDA)

- Question:
 - “What are the characteristics which have the most impact on the car price?”
- Preliminary step in data analysis to:
 - Summarize main characteristics of the data
 - Gain better understanding of the data set
 - Uncover relationships between variables
 - Extract important variables

Learning Objectives:

- Descriptive Statistics
- GroupBy
- Correlation
- Correlation - Statistics

Descriptive Statistics - Describe()

- Summarize statistics using pandas `describe()` method
 - `df.describe()`
- Summarize categorical data is by using the `value_counts()` method
- Box Plot
- Scatter Plot
 - each observation represented as a point
 - scatter plot show the relationship between two variables
 - predictor/independent variables on x-axis
 - target/dependent variables on y-axis

Grouping data

`groupby`

- use `df.groupby()` method:
 - can be applied on categorical variables
 - group data into categories
 - single or multiple variables

Groupby()- Example

```
df_test = df[['drive-wheels', 'body-style', 'price']]
df_grp = df_test.groupby(['drive-wheels', 'body-style'], as_index=False).mean()
df_grp
```



	drive-wheels	body-style	price
0	4wd	convertible	20239.229524
1	4wd	sedan	12647.333333
2	4wd	wagon	9095.750000
3	fwd	convertible	11595.000000
4	fwd	hardtop	8249.000000
5	fwd	hatchback	8396.387755
6	fwd	sedan	9811.800000
7	fwd	wagon	9997.333333
8	rwd	convertible	23949.600000
9	rwd	hardtop	24202.714286
10	rwd	hatchback	14337.777778
11	rwd	sedan	21711.833333

A table of this form isn't the easiest to read and also not very easy to visualize.

To make it easier to understand, we can transform this table to a pivot table by using the `pivot` method.

`pivot`

Pandas method - Pivot()

- One variable displayed along the columns and the other variable displayed along the rows.

```
df_pivot = df_grp.pivot(index= 'drive-wheels', columns='body-style')
```

	price				
body-style	convertible	hardtop	hatchback	sedan	wagon
drive-wheels					
4wd	20239.229524	20239.229524	7603.000000	12647.333333	9095.750000
fwd	11595.000000	8249.000000	8396.387755	9811.800000	9997.333333
rwd	23949.600000	24202.714286	14337.777778	21711.833333	16994.222222

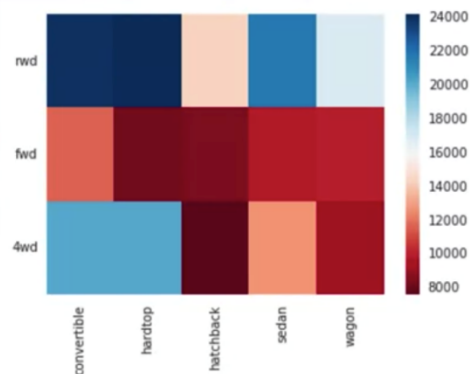
The price data now becomes a rectangular grid, which is easier to visualize. This is similar to what is usually done in Excel **spreadsheets**. Another way to represent the pivot table is using a **heat map** plot.

Heatmap

Heatmap

- Plot target variable over multiple variables

```
plt.pcolor(df_pivot, cmap='RdBu')  
plt.colorbar()  
plt.show()
```

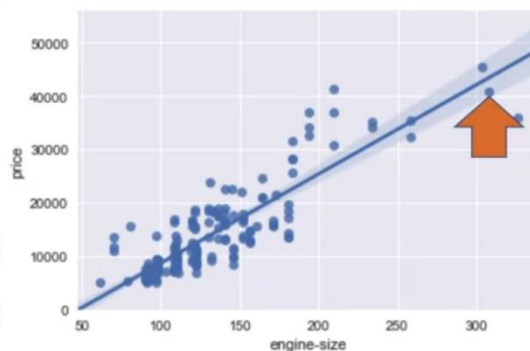


Correlation

Correlation - Positive Linear Relationship

- Correlation between two features (engine-size and price).

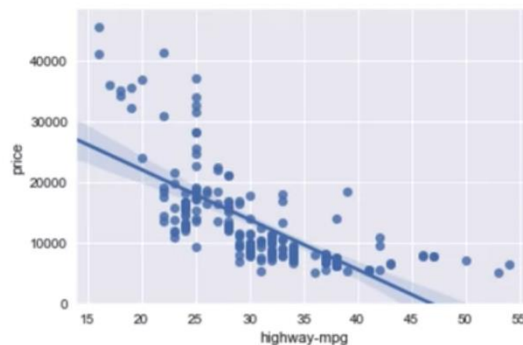
```
sns.regplot(x="engine-size", y="price", data=df)  
plt.ylim(0,)
```



Correlation - Negative Linear Relationship

- Correlation between two features (highway-mpg and price).

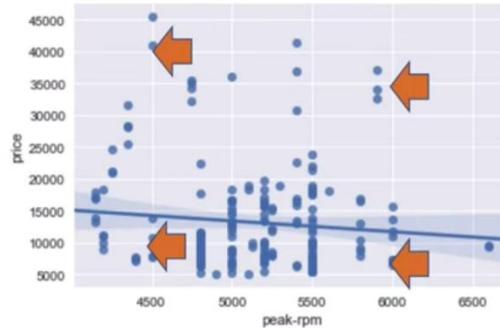
```
sns.regplot(x="highway-mpg", y="price", data=df)  
plt.ylim(0,)
```



Correlation - Negative Linear Relationship

- Weak correlation between two features (peak-rpm and price).

```
sns.regplot(x="peak-rpm", y="price", data=df)  
plt.ylim(0,)
```

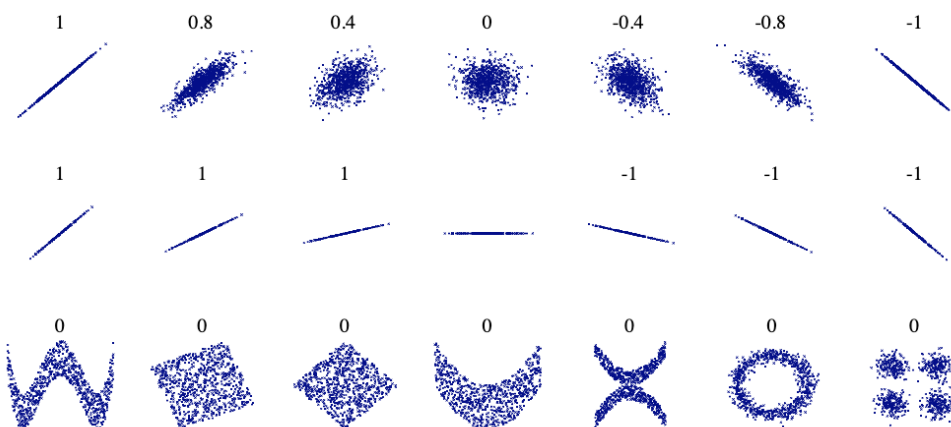


Correlation - Statistics

Pearson Correlation

Pearson Correlation

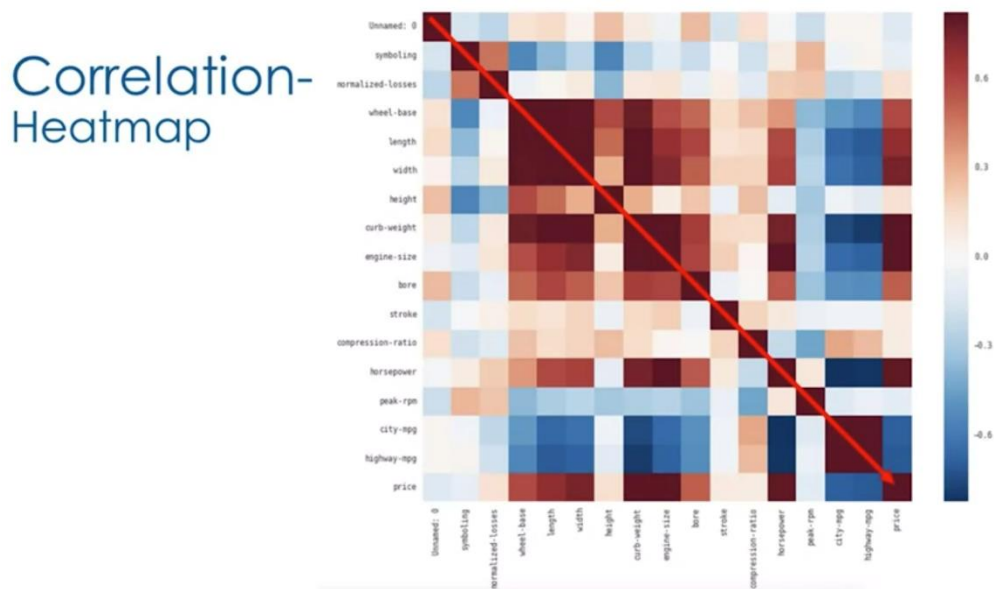
- Measure the strength of the correlation between two features.
 - Correlation coefficient
 - P-value
- Correlation coefficient
 - Close to +1: Large Positive relationship
 - Close to -1: Large Negative relationship
 - Close to 0: No relationship
- P-value
 - P-value < 0.001 **Strong** certainty in the result
 - P-value < 0.05 **Moderate** certainty in the result
 - P-value < 0.1 **Weak** certainty in the result
 - P-value > 0.1 **No** certainty in the result
- Strong Correlation:
 - Correlation coefficient close to 1 or -1
 - P value less than 0.001



The correlation reflects the noisiness and direction of a linear relationship (top row), but not the slope of that relationship (middle), nor many aspects of nonlinear relationships (bottom). N.B.: the

figure in the center has a slope of 0 but in that case the correlation coefficient is undefined because the variance of Y is zero.

Correlation Heatmap



Association between two categorical variables: Chi-Square

Categorical variables

- use the Chi-square Test for Association (denoted as χ^2)
- The test is intended to test how likely it is that an observed distribution is due to chance

Chi-Square Test of association

- The Chi-square tests a null hypothesis that the variables are independent.
- The Chi-square does not tell you the type of relationship that exists between both variables; but only that a relationship exists.

See also: [Chi-Square Test of Independence](#)

Categorical variables

- Is there an association between fuel-type and aspiration?

$$\chi^2 = \sum_{i=1}^n \frac{(O_i - E_i)^2}{E_i}$$

Observed value

	Standard	Turbo	Total
diesel	7	13	20
gas	161	24	185
Total	168	37	205

Row total * Column total

Grand total

Categorical variables

$$\chi^2 = \sum_{i=1}^n \frac{(O_i - E_i)^2}{E_i}$$

Expected value

	Standard	Turbo	Total
diesel	7	13	20
gas	161	24	185
Total	168	37	205

Observed value

Aspiration Fuel-type	Standard	Turbo	Total
Diesel	16.39	3.61	20
Gas	151.61	33.39	185
Total	168	37	

Row total * Column total

Grand total

Categorical variables

$$\chi^2 = \sum_{i=1}^n \frac{(O_i - E_i)^2}{E_i}$$

Degree of freedom = (row-1)*(column-1)

$\chi^2 = 29.6$

Percentage Points of the Chi-Square Distribution

Degrees of Freedom	0.99	0.95	0.90	0.75	0.50	0.25	0.10	0.05
1	0.000	0.004	0.016	0.102	0.455	1.32	2.71	3.84
2	0.020	0.103	0.211	0.575	1.386	2.77	4.61	5.99
3	0.115	0.352	0.584	1.212	2.366	4.11	6.25	7.81
4	0.297	0.711	1.064	1.923	3.357	5.39	7.78	9.49
5	0.554	1.145	1.610	2.675	4.351	6.63	9.24	11.07
6	0.872	1.635	2.204	3.455	5.348	7.84	10.64	12.59
7	1.239	2.167	2.833	4.255	6.346	9.04	12.02	14.07
8	1.647	2.733	3.490	5.071	7.344	10.22	13.36	15.51
9	2.088	3.325	4.168	5.899	8.343	11.39	14.68	16.92

P-value < 0.05, we reject the null hypothesis that the two variables are independent and conclude that there is evidence of association between fuel-type and aspiration.

Categorical variables

$$\chi^2 = \sum_{i=1}^n \frac{(O_i - E_i)^2}{E_i}$$

```
scipy.stats.chi2_contingency(cont_table, correction = True)
```

```
(29.605759385109046,  
5.2947382636786724e-08,  
1,  
array([[ 16.3902439,  3.6097561],  
       [151.6097561, 33.3902439]]))
```

	Standard	Turbo	Total
diesel	7	13	20
gas	161	24	185
Total	168	37	205

P-value of < 0.05, we reject the null hypothesis that the two variables are independent and conclude that there is evidence of association between fuel-type and aspiration.

Model Development

- simple linear regression
- multiple linear regression
- polynomial regression

Linear Regression and Multiple Linear Regression

Fitting a Simple Linear Model

- We define the predictor variable and target variable

```
X = df[['highway-mpg']]  
Y = df['price']
```

- Then use `lm.fit(X, Y)` to fit the model, i.e. find the parameters b_0 and b_1

```
lm.fit(X, Y)
```

- We can obtain a prediction

```
Yhat=lm.predict(X)
```

Yhat	X
2	5
:	
3	4

SLR – Estimated Linear Model

- We can view the intercept (b_0): `lm.intercept_`
38423.305858
- We can also view the slope (b_1): `lm.coef_`
-821.73337832
- The Relationship between Price and Highway MPG is given by:
- **Price = 38423.31 - 821.73 * highway-mpg**

$$\hat{Y} = b_0 + b_1x$$

Fitting a Multiple Linear Model Estimator

1. We can extract the for 4 predictor variables and store them in the variable Z

```
Z = df[['horsepower', 'curb-weight', 'engine-size', 'highway-mpg']]
```

2. Then train the model as before:

```
lm.fit(Z, df['price'])
```

3. We can also obtain a prediction

```
Yhat=lm.predict(X)
```

x_1	x_2	x_3	x_4		Yhat
3	5	-4	3	→	2
:	:	:	:		:
2	4	2	-4		3

MLR – Estimated Linear Model

1. Find the intercept (b_0)

```
lm.intercept_  
-15678.742628061467
```

2. Find the coefficients (b_1, b_2, b_3, b_4)

```
lm.coef_  
array([52.65851272, 4.69878948, 81.95906216, 33.58258185])
```

The Estimated Linear Model:

- **Price = -15678.74 + (52.66) * horsepower + (4.70) * curb-weight + (81.96) * engine-size + (33.58) * highway-mpg**

Model Evaluation using Visualization

Regression Plot

Regression plot gives us a good estimate of:

- the relationship between two variables
- the strength of the correlation
- the direction of the relationship (positive or negative)

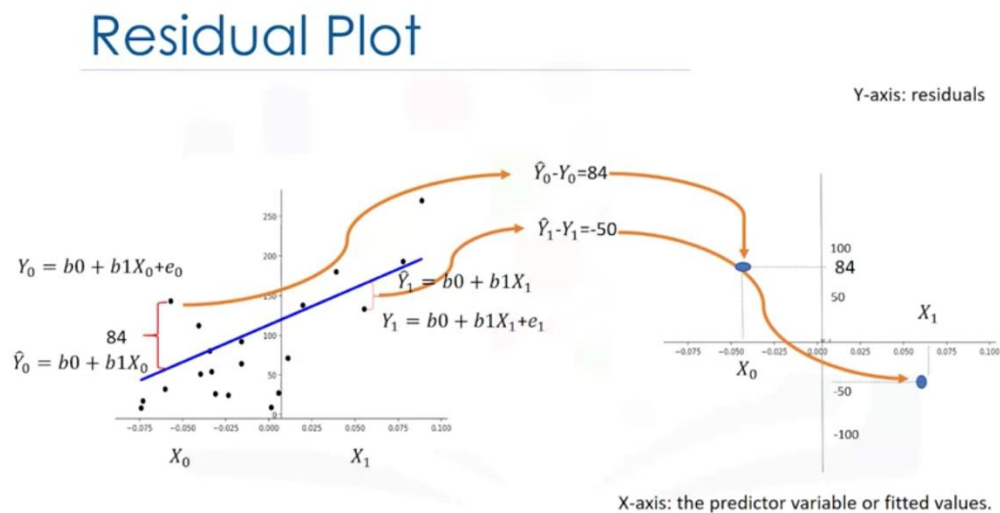
Regression plot shows us a combination of:

- the scatterplot: where each point represents a different y
- the fitted linear regression line (\hat{y})

```
import seaborn as sns

sns.regplot(x="highway-mpg", y="price", data=df)
plt.ylim(0,)
```

Residual Plot



We expect to see the results to have **zero mean**, distributed **evenly** around the x axis with similar variance.

```
import seaborn as sns

sns.residplot(df["highway-mpg"], df["price"])
```

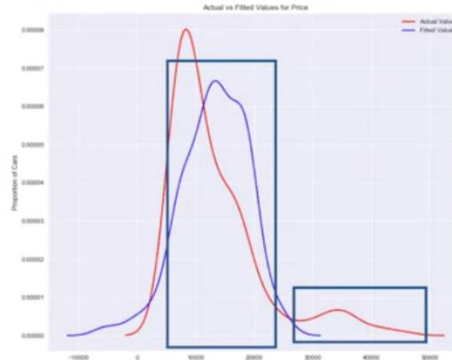
Distribution Plots

A distribution plot counts the predicted value versus the actual value. These plots are extremely useful for visualizing models with more than one independent variable or feature.

Distribution Plots

Compare the distribution plots:

- The fitted values that result from the model
- The actual values



```
import seaborn as sns

ax1 = sns.distplot(df["price"], hist=False, color="r", label="Actual Value")

sns.distplot(Yhat, hist=False, color="b", label="Fitted Value", ax=ax1)
```

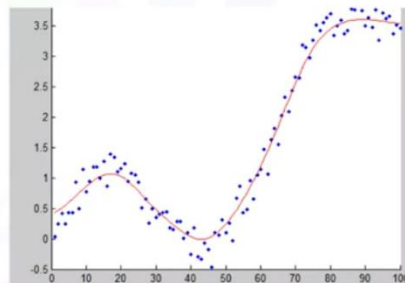
Polynomial Regression and Pipelines

Polynomial Regressions

- A special case of the general linear regression model
- Useful for describing curvilinear relationships

Curvilinear relationships:

By squaring or setting higher-order terms of the predictor variables



Polynomial Regression

- Quadratic – 2nd order

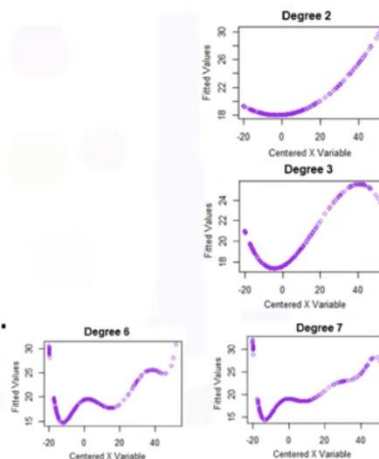
$$\hat{Y} = b_0 + b_1 x_1 + b_2 (x_1)^2$$

- Cubic – 3rd order

$$\hat{Y} = b_0 + b_1 x_1 + b_2 (x_1)^2 + b_3 (x_1)^3$$

- Higher order

$$\hat{Y} = b_0 + b_1 x_1 + b_2 (x_1)^2 + b_3 (x_1)^3 + \dots$$



Polynomial Regression

1. Calculate Polynomial of 3rd order

```
f=np.polyfit(x,y,3)
```

```
p=np.polyld(f)
```

2. We can print out the model

```
print (p)
```

$$-1.557(x_1)^3 + 204.8(x_1)^2 + 8965 x_1 + 1.37 \times 10^5$$

Polynomial Regression with More than One Dimension

- We can also have multi dimensional polynomial linear regression

$$\hat{Y} = b_0 + b_1 X_1 + b_2 X_2 + b_3 X_1 X_2 + b_4 (X_1)^2 + b_5 (X_2)^2 + \dots$$

Numpy's polyfit function cannot perform this type of regression. We use the preprocessing library in scikit-learn to create a polynomial feature object.

```
from sklearn.preprocessing import PolynomialFeatures

pr = PolynomialFeatures(degree=2, include_bias=False)
x_poly = pr.fit_transform(x[['horsepower', 'curb-weight']])
```


Polynomial Regression with More than One Dimension

```
pr=PolynomialFeatures(degree=2)
```

X_1	X_2
1	2

```
pr=PolynomialFeatures(degree=2,include_bias=False)  
pr.fit_transform([[1,2]])
```



X_1	X_2	X_1X_2	X_1^2	X_2^2
1	2	(1) 2	1	(2) ²

1	2	2	1	4
---	---	---	---	---

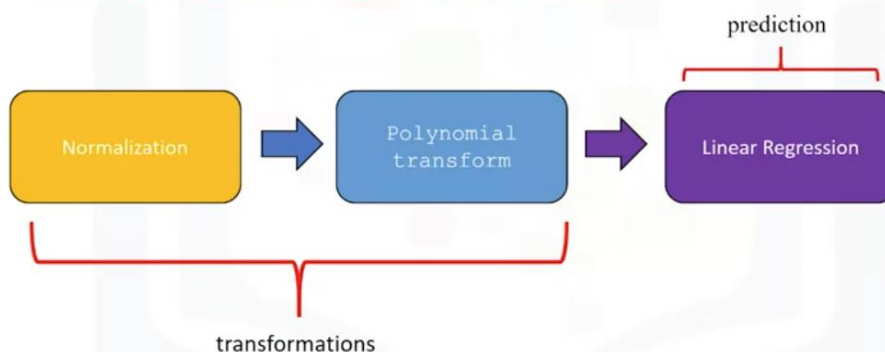
As the dimension of the data gets larger, we may want to normalize multiple features in scikit-learn. Instead we can use the preprocessing module to simplify many tasks. For example, we can standardize each feature simultaneously. We import `StandardScaler`.

```
from sklearn.preprocessing import StandardScaler  
SCALE = StandardScaler()  
SCALE.fit(x_data[['horsepower', 'highway-mpg']])  
x_scale = SCALE.transform(x_data[['horsepower', 'highway-mpg']])
```

We can simplify our code by using a pipeline library.

Pipelines

- There are many steps to getting a prediction



```
from sklearn.preprocessing import PolynomialFeatures  
from sklearn.preprocessing import StandardScaler  
from sklearn.linear_model import LinearRegression  
from sklearn.pipeline import Pipeline  
  
Input = [('scale', StandardScaler()), ('polynomial', PolynomialFeatures(degree=2),...),  
('model', LinearRegression())]  
pipe = Pipeline(Input)  
pipe.fit(df[['horsepower', 'curb-weight', 'engine-size', 'highway-mpg']], y)  
yhat = pipe.predict(X[['horsepower', 'curb-weight', 'engine-size', 'highway-mpg']])
```

Measures for In-Sample Evaluation

Measures for In-Sample Evaluation

- A way to numerically determine how good the model fits on dataset
- Two important measures to determine the fit of a model:
 - Mean Squared Error (MSE)
 - R-squared (R^2)

Mean Squared Error (MSE)

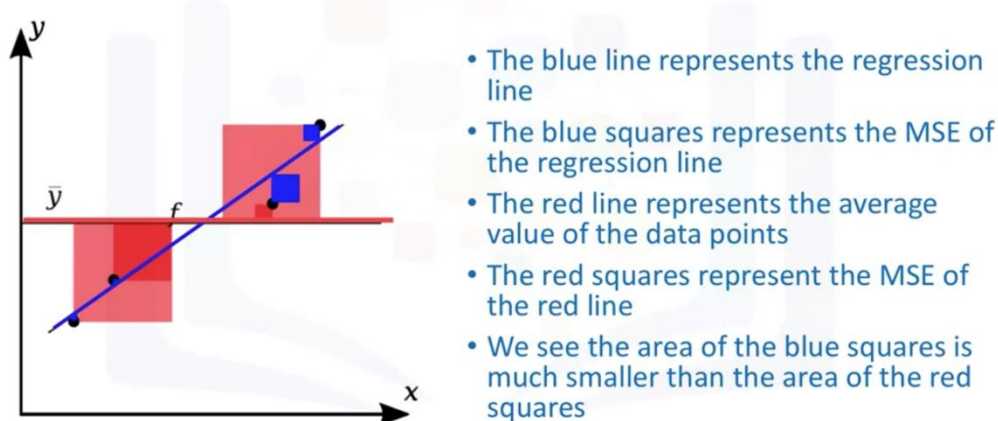
```
from sklearn.metrics import mean_square_error  
  
mean_square_error(df['price'], Y_predict_simple_fit)
```

R-squared

- The [Coefficient of Determination](#) or R-squared (R^2)
- Is a measure to determine how close the data is to the fitted regression line.
- R^2 : the percentage of variation of the target variable (Y) that is explained by the linear model.
- think about as comparing a regression model to a simple model i.e. the mean of the data points

$$R^2 = (1 - (\text{MSE of regression line}) / (\text{MSE of the average of the data}))$$

Coefficient of Determination (R^2)



- Generally the values of the MSE are between 0 and 1
- We can calculate the R^2 as follows

```
X = df[['highway-mpg']]
Y = df['price']
lm.fit(X, Y)
lm.score(X, Y) # 0.496591188
```

We can say that approximately **49.695%** of the variation of price is explained by this simple linear model.

Comparing MLR and SLR

Does a lower Mean Square Error imply better fit?

- Not necessarily
- 1. Mean Square Error for a Multiple Linear Regression Model will be smaller than the Mean Square Error for a Simple Linear Regression model, since the errors of the data will decrease when more variables are included in the model
- 2. Polynomial regression will also have a smaller Mean Square Error than the linear regular regression
- 3. In the next section we will look at more accurate ways to evaluate the model

Model Evaluation and Refinement

Training/Testing Sets

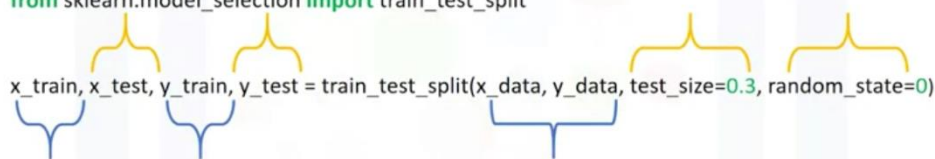
- Split dataset into:
 - Training set (70%)
 - Testing set (30%)
- Build and train the model with a training set
- Use testing set to assess the performance of a predictive model
- When we have completed testing our model we should use all the data to train the model to get the best performance

Function `train_test_split()`

- Split data into random train and test subsets

```
from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split(x_data, y_data, test_size=0.3, random_state=0)
```



- `x_data`: features or independent variables
- `y_data`: dataset target: `df['price']`
- `x_train`, `y_train`: parts of available data as training set
- `x_test`, `y_test`: parts of available data as testing set
- `test_size`: percentage of the data for testing (here 30%)
- `random_state`: number generator used for random sampling

Function `cross_val_score()`

One of the most common out of sample evaluation metrics is [cross-validation](#).

- In this method, the dataset is split into K equal groups.
- Each group is referred to as a fold. For example, four folds. Some of the folds can be used as a training set which we use to train the model and the remaining parts are used as a test set, which we use to test the model.
- For example, we can use three folds for training, then use one fold for testing. This is repeated until each partition is used for both training and testing.
- At the end, we use the average results as the estimate of out-of-sample error.
- The evaluation metric depends on the model, for example, the r squared.

The simplest way to apply cross-validation is to call the `cross_val_score` function, which performs multiple out-of-sample evaluations.

```
from sklearn.model_selection import cross_val_score

score = cross_val_score(lr, x_data, y_data, cv=3)
np.mean(scores)
```

Function `cross_val_predict()`

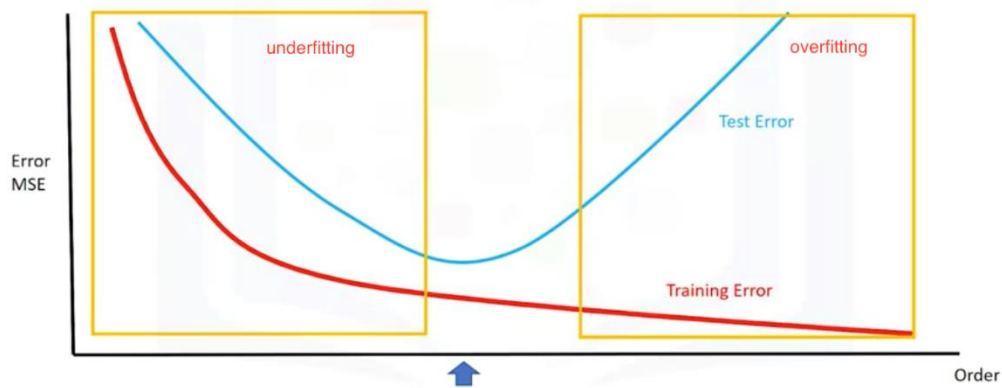
- It returns the prediction that was obtained for each element when it was in the test set
- Has a similar interface to `cross_val_score()`

```
from sklearn.model_selection import cross_val_predict

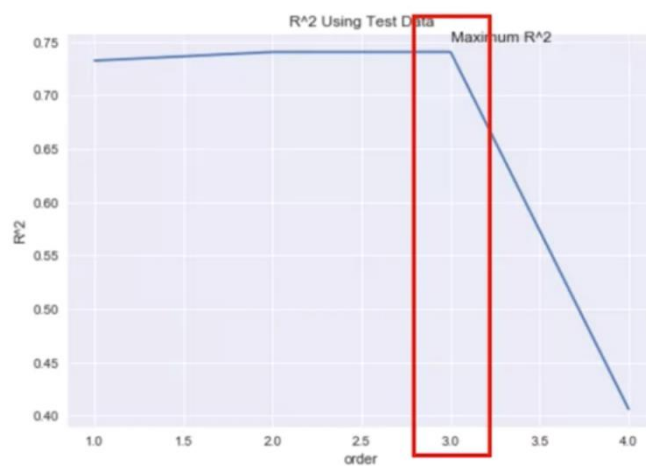
yhat = cross_val_predict(lr2e, x_data, y_data, cv=3)
```

Overfitting, Underfitting and Model Selection

Model Selection



Model Selection



Calculate different R-squared values as follows:

```
Rsqu_test = []
order = [1,2,3,4]

for n in order:
    pr = PolynomialFeatures(degree=n)
    x_train_pr = pr.fit_transform(x_train[['horsepower']])
    x_test_pr = pr.fit_transform(x_test[['horsepower']])
    lr.fit(x_train_pr, y_train)
    Rsqu_test.append(lr.score(x_test_pr, y_test))
```

Ridge Regression

Ridge regression is a regression that is employed in a Multiple regression model when Multicollinearity occurs. Multicollinearity is when there is a strong relationship among the independent variables. Ridge regression is very common with polynomial regression.

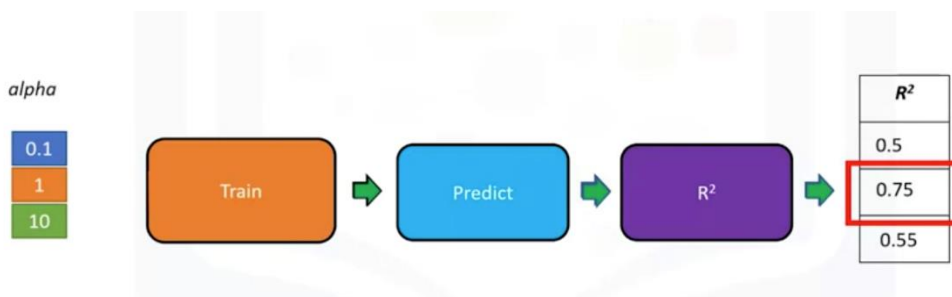
Ridge Regression

$$\hat{y} = 1 + 2x - 3x^2 - 2x^3 - 12x^4 - 40x^5 + 80x^6 + 71x^7 - 141x^8 - 38x^9 + 75x^{10}$$

Alpha	x	x ²	x ³	x ⁴	x ⁵	x ⁶	x ⁷	x ⁸	x ⁹	x ¹⁰
0	2	-3	-2	-12	-40	80	71	-141	-38	75
0.001	2	-3	-7	5	4	-6	4	-4	4	6
0.01	1	-2	-5	-0.04	0.15	-1	1	-0.5	0.3	1
1	0.5	-1	-1	-0.614	0.70	-0.38	-0.56	-0.21	-0.5	-0.1
10	0	-0.5	-0.3	-0.37	-0.30	-0.30	-0.22	-0.22	-0.22	-0.17

The column corresponds to the different polynomial coefficients, and the rows correspond to the different values of alpha.

- As alpha increases, the parameters get smaller. This is most evident for the higher order polynomial features.
- But Alpha must be selected carefully.
 - If alpha is too large, the coefficients will approach zero and underfit the data.
 - If alpha is zero, the overfitting is evident.



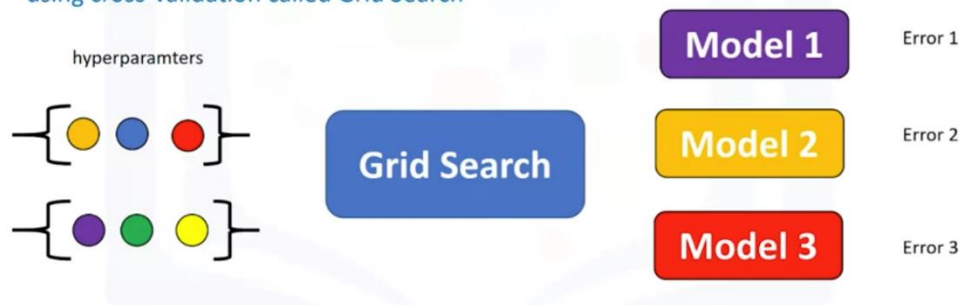
Grid Search

- The term alpha in Ridge regression is called a **hyperparameter**.
- Scikit-learn has a means of automatically iterating over these hyperparameters using cross-validation called **Grid Search**.

[Grid Search](#) takes the model or objects you would like to train and different values of the hyperparameters. It then calculates the mean square error or R-squared for various hyperparameter values, allowing you to choose the best values.

Hyperparameters

- In the last section, the term alpha in Ridge regression is called a hyperparameter
- Scikit-learn has a means of automatically iterating over these hyperparameters using cross-validation called Grid Search



Use the validation dataset to pick the best hyperparameters.

```
from sklearn.linear_model import Ridge
from sklearn.model_selection import GridSearchCV

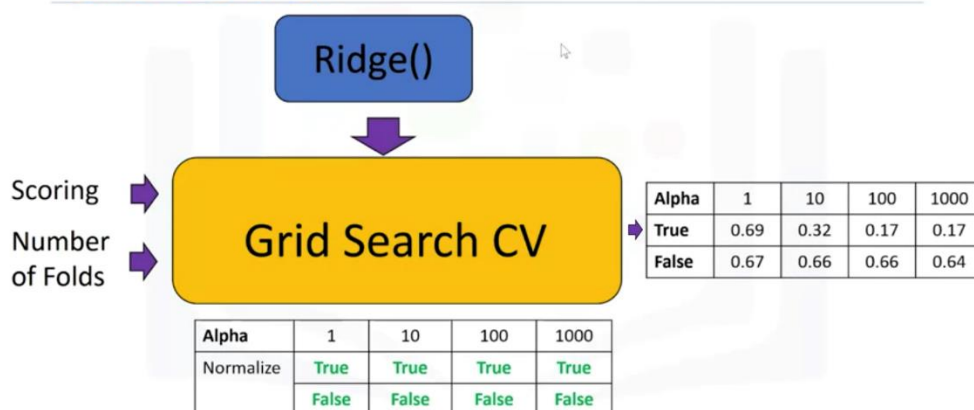
parameters1 = [{'alpha': [0.001, 0.1, 1, 10, 100, 1000, 10000, 100000]}]

RR = Ridge()
Grid1 = GridSearchCV(RR, parameters1, cv=4)
Grid1.fit(x_data[['horsepower', 'curb-weight', 'engine-size', 'highway-mpg']], y_data)
Grid1.best_estimator_

scores = Grid1.cv_results_
scores['mean_test_score']
```

What are the advantages of Grid Search is how quickly we can test **multiple parameters**.

Grid Search



```
from sklearn.linear_model import Ridge
from sklearn.model_selection import GridSearchCV

parameters2 = [{'alpha': [0.001, 0.1, 1, 10, 100], 'normalize': [True, False]}]

RR = Ridge()
Grid1 = GridSearchCV(RR, parameters2, cv=4)
```



```
Grid1.fit(x_data[['horsepower', 'curb-weight', 'engine-size', 'highway-mpg']], y_data)
Grid1.best_estimator_
```

```
scores = Grid1.cv_results_
```

```
for param, mean_val, mean_test in zip(scores['params'], scores['mean_test_score'],
scores['mean_train_score']):
    print(param, "R^2 on test data:", mean_val, "R^2 on train data:", mean_test)
```

```
{'alpha': 0.001, 'normalize': True} R^2 on tesst data: 0.66605547293 R^2 on train data: 0.814001968709
{'alpha': 0.001, 'normalize': False} R^2 on tesst data: 0.665488366584 R^2 on train data: 0.814002698797
{'alpha': 0.1, 'normalize': True} R^2 on tesst data: 0.694175625356 R^2 on train data: 0.810546768311
{'alpha': 0.1, 'normalize': False} R^2 on tesst data: 0.665488937796 R^2 on train data: 0.814002698794
{'alpha': 1, 'normalize': True} R^2 on tesst data: 0.690486934584 R^2 on train data: 0.749104440368
{'alpha': 1, 'normalize': False} R^2 on tesst data: 0.665494127178 R^2 on train data: 0.814002698472
{'alpha': 10, 'normalize': True} R^2 on tesst data: 0.321376875232 R^2 on train data: 0.341856042902
{'alpha': 10, 'normalize': False} R^2 on tesst data: 0.665545680812 R^2 on train data: 0.8140026666
{'alpha': 100, 'normalize': True} R^2 on tesst data: 0.0170551710263 R^2 on train data: 0.0496044796826
{'alpha': 100, 'normalize': False} R^2 on tesst data: 0.666029359996 R^2 on train data: 0.813999791851
{'alpha': 1000, 'normalize': True} R^2 on tesst data: -0.0301961745066 R^2 on train data: 0.005184451599
{'alpha': 1000, 'normalize': False} R^2 on tesst data: 0.668968215369 R^2 on train data: 0.813870488264
{'alpha': 10000, 'normalize': True} R^2 on tesst data: -0.0351687400461 R^2 on train data: 0.000520784757979
{'alpha': 10000, 'normalize': False} R^2 on tesst data: 0.673346359342 R^2 on train data: 0.812583743226
{'alpha': 100000, 'normalize': True} R^2 on tesst data: -0.0356685844558 R^2 on train data: 5.2101975528e-05
{'alpha': 100000, 'normalize': False} R^2 on tesst data: 0.657818838432 R^2 on train data: 0.789541446486
{'alpha': 100000, 'normalize': True} R^2 on tesst data: -0.0356685844558 R^2 on train data: 5.2101975528e-05
{'alpha': 100000, 'normalize': False} R^2 on tesst data: 0.657818838432 R^2 on train data: 0.789541446486
```